

# OpenMP

---

Martin Kruliš  
Jakub Yaghob





# About

- OpenMP (<http://www.openmp.org>)
  - API for multi-threaded, shared memory parallelism
  - Supported directly by compilers
    - **C/C++** and Fortran
    - Activated by compiler directive (e.g., `g++ -fopenmp`)
  - Three components
    - Compiler directives (pragmas)
    - Runtime library resources (functions)
    - Environment variables (defaults, runtime configuration)

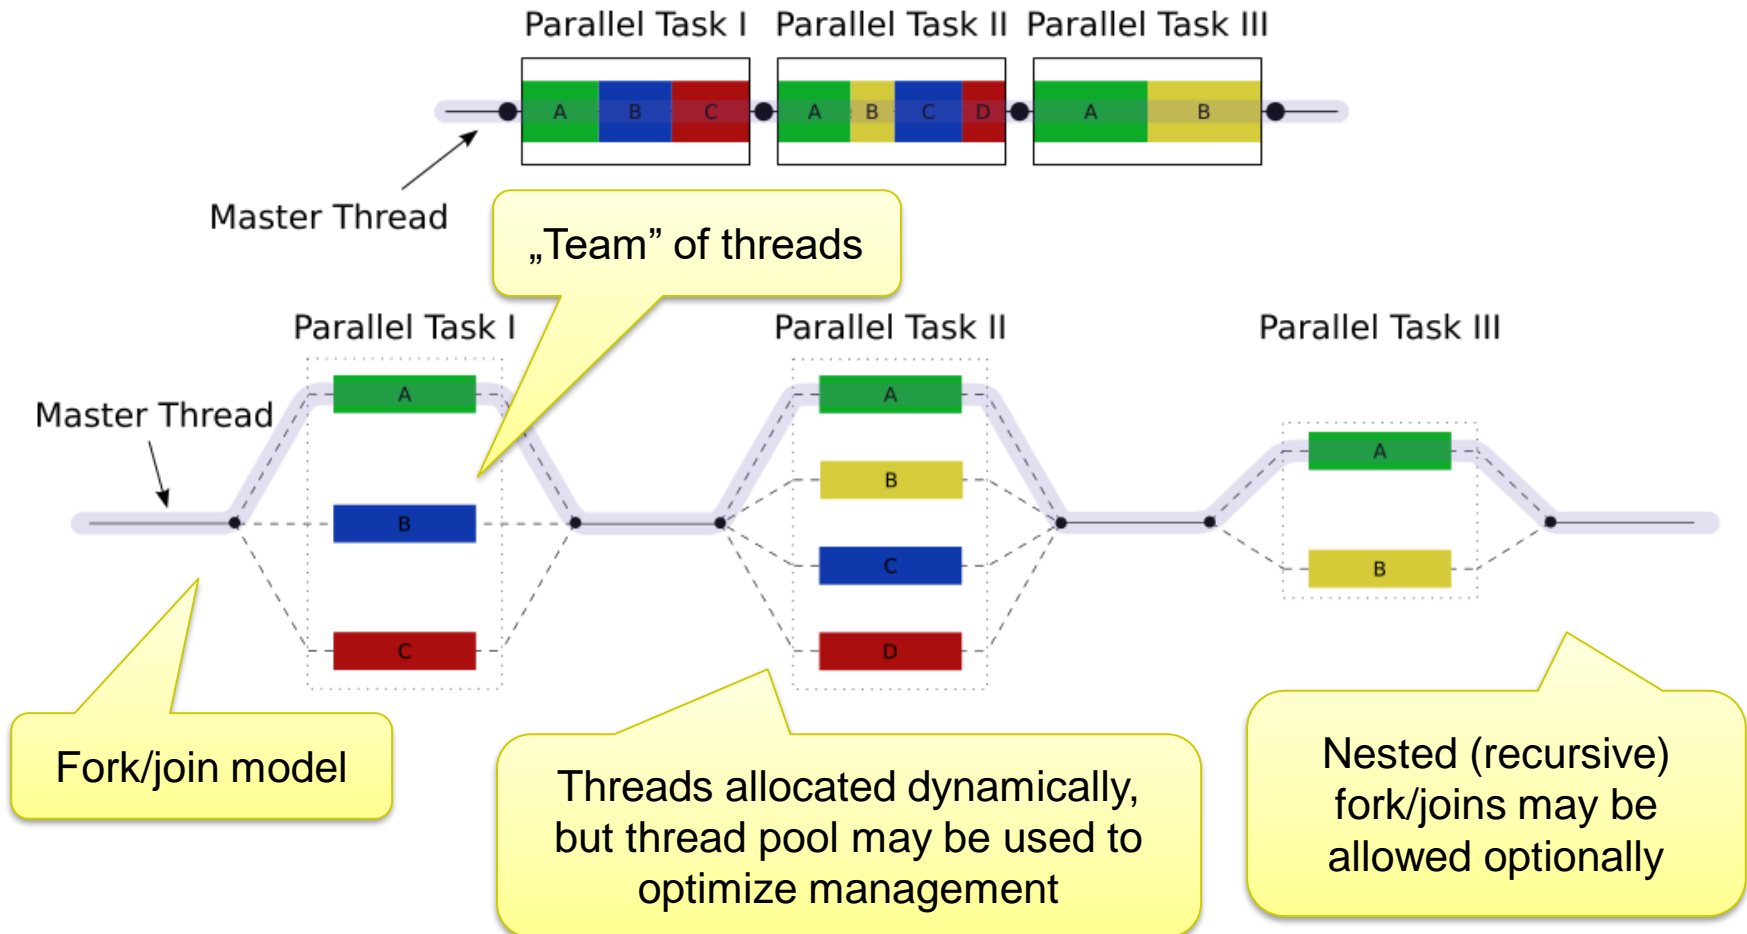


# History

- Specification versions
  - 1.0 – C/C++ and FORTRAN versions (1997-1998)
  - 2.0 – C/C++ and FORTRAN versions (2000-2002)
    - Supported by MSVC
  - 2.5 – combined C/C++ and FORTRAN (2005)
  - 3.0 – tasks (2008)
  - 3.1 – additional extensions for tasks and atomics (2011)
  - 4.0 – SIMD, cancellation, array sections (2013)
  - 4.5 – Fortran 2003 (2015)
    - Widely available (gcc 6+, Intel compilers, clang)
  - 5.0 – memory model, accelerators (GPU), unified shared memory, iterators, debugger support (2018)
  - 5.1 – full C++20 and Fortran 2008 support, C++ attributes (2020)
    - Partial support (gcc 11, Intel compilers, clang)
  - 5.2 – improvements, refinements
  - 6.0 – should be released 2024



# Execution Model





# Memory model

- Relaxed consistency, shared memory model
- Memory
  - Storage location may be associated with one or more devices
  - Only threads on devices may access it
- Thread private memory
  - Not accessible by other threads
- Load/store are not guaranteed to be atomic
- Memory operations are considered unordered, unless some defined cases



# Pragmas

- Basic Syntax

**#pragma omp** *directive* [*clause*, ...]

- Code should work without the pragmas (as serial)
- Pragmas may be used to
  - Spawn a parallel region
  - Divide workload among threads
  - Serialize sections of code
  - Synchronization



# Directives as C++ attributes

- Effort to remove C++ preprocessor

```
[[ omp :: directive( directive [clause, ...]  
) ]]
```

or

```
[[ using omp : directive( directive  
[clause, ...] ) ]]
```

- Attribute directives that apply to the same statement are unordered
- Ordering can be imposed

```
[[ omp :: sequence( [omp::]directive ,  
[omp::]directive ) ]]
```

# Directives as C++ attributes – demo



```
[[ omp::sequence(  
directive(parallel) ,  
directive(for)) ]]
```

```
for(...) {}
```

```
#pragma omp parallel
```

```
#pragma omp for
```

```
for(...) {}
```





# Parallel Region

- Spawning a thread team

```
#pragma omp parallel
```

Spawns one thread per CPU code by default

```
{
```

```
    int tid = omp_get_thread_num() ;
```

Thread index within its team (master == 0)

```
    // use tid to do thread's bidding
```

```
}
```

- Implicit barrier at the end
- No branching/goto-s that will cause the program to jump in/out to/of parallel blocks
- Regular branching or function calls are OK



# Variable Scope

- Variables Scope
  - Private – a copy per each thread
    - Local variables in a parallel block are implicitly private
  - Shared – all threads share the same variable
    - Synchronization may be required

```
int x, y, id;  
#pragma omp parallel private(id), shared(x,y)  
{  
    id = x + omp_get_thread_num();  
    ...  
}
```






# Variable Scope

- Private Scope
  - Most variables are private by default
    - Variables declared inside the parallel block
    - All non-static variables in called functions
    - Values are not initialized (at the beginning and the end of the block)
      - Except for classes (default constructor must be accessible)
- Other scopes and additional clauses
  - Will be presented later



# Work-sharing Constructs

- Divide the work in parallel block





# Work-sharing Constructs

- For-loop

```
#pragma omp parallel
```

Spawns the threads

```
{
```

```
    #pragma omp for
```

Divide for-loop workload among the threads

```
    for (int i = 0; i < N; ++i) { ... }
```

```
}
```

- Additional clauses of **for** directive

- **schedule(strategy)** – scheduling (work division)
  - static, dynamic, guided, runtime, auto
- **collapse(n)** – encompass nested for-loops



# Work-sharing Constructs

- For-loop
  - `#pragma omp ordered`
    - A block inside for-loop that must be executed in exactly the same order, as if the code was serial
  - `#pragma omp parallel for`
    - Shorthand for both creating parallel block and apply it on a parallel for-loop
    - May have clauses applied to both `parallel` and `for` directives
    - Probably the most often used construct in OpenMP



# Work-sharing Constructs

- For-loop Pitfalls

- Use `#pragma omp for` outside of `#pragma omp parallel` block
  - Has no effect, there is only one thread available!
- Forgetting the `for` itself

```
#pragma omp parallel
for (int i = 0; i < N; ++i) { ... }
```

  - The entire loop is executed by **ALL** threads



# Work-sharing Constructs

- Sections

- Independent blocks of code executed concurrently

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        load_player_data();
        #pragma omp section
        load_game_maps();
        ...
    }
}
```





# Work-sharing Constructs

- Single

- Code executed by single thread from group

```
#pragma omp parallel
```

```
{
```

```
    #pragma omp for
```

```
    for (...) ...
```

```
    #pragma omp single
```

```
    report_progress();
```

Only one thread reports the progress

```
    #pragma omp for
```

```
    ...
```

```
}
```



# Work-sharing Constructs

- Synchronization
  - Implicit barrier at the end of each construct
    - for, sections, single
  - **nowait** clause
    - Removes the barrier



# Synchronization Constructs

- Synchronization Directives
  - To be used within a parallel block
  - **#pragma omp master**
    - Region being executed only by the master thread
      - Similar to **#pragma omp single**
  - **#pragma omp critical** [*name*]
    - Standard critical section with lock guard
      - Only one thread may be in the section at a time
    - If a name is provided, all sections with the same name are interconnected (use the same lock)



# Synchronization Constructs

- Synchronization Directives
  - `#pragma omp barrier`
    - All threads in a team must meet on a barrier
  - `#pragma omp atomic`
    - Followed by a statement like `x += expr;` or `++x;`
    - Allowed operations: `+`, `*`, `-`, `/`, `&`, `^`, `|`, `<<`, or `>>`
  - `#pragma omp flush` [*memory-ordering*]
    - Make sure changes of shared variables become visible
    - Executed implicitly for many directives
      - barrier, critical, parallel, ...



# Tasks

- Tasks
  - Pieces of code that may be executed by a different thread (within a parallel section)  
`#pragma omp task`
  - Additional clauses
    - `untied` – different thread may resume task after yield
    - `priority(p)` – scheduling hint
    - `depend(list)` – tasks that must conclude first
    - And few other that control when and who can execute the task



# Tasks

- Tasks Synchronization
  - `#pragma omp taskwait`
    - Wait for all child tasks (spawned in this task)
  - `#pragma omp taskyield`
    - Placed as statement inside a task
    - The processing thread may suspend this task and pick up another task
  - `#pragma omp taskgroup`
    - Spawning thread will not continue unless all tasks in the group are completed



# Data Sharing Management

- Data-related Clauses
  - We already know **private** and **shared**
  - **firstprivate** – similar to private, but the value is initialized using the value of the main thread
  - **lastprivate** – value after parallel block is set to the last value, that would be computed in serial processing
  - **copyprivate** – used with single block, last value is broadcasted to all threads



# Data Sharing Management

- Reduction

- Variable is private and reduced at the end
- **#pragma omp ... reduction(*op:list*)**

- Op represents operation (+, \*, &, |, ...)
  - Each operation has its own default (e.g., + has 0)
- List of variables

- Custom reducers

```
#pragma omp declare reduction (identifier :  
typename-list : combiner) [initializer-clause]
```





# Synchronization Constructs

- Local Thread Storage

**#pragma omp threadprivate**(*list*)

- Variables are made private for each thread
  - No connection to explicit parallel block
  - Values persist between blocks
- Variables must be either global or static
- **copyin**(*list*) clause (of parallel block)
  - Similar to **firstprivate**
  - Threadprivate variables are initialized by master value



# Runtime Library

- Runtime Library
  - `#include <omp.h>` usually required
- Functions
  - `void omp_set_num_threads(int threads)`
    - Set # of threads in next parallel region
  - `int omp_get_num_threads(void)`
    - Get # of threads in current team
  - `int omp_get_max_threads(void)`
    - Current maximum of threads in a parallel region



# Runtime Library

- Functions
  - `int omp_get_thread_num(void)`
    - Current thread index within a team (master == 0)
  - `int omp_get_num_procs(void)`
    - Actual number of CPU cores (available)
  - `int omp_in_parallel(void)`
    - True, if the code is executed in parallel
  - `omp_set_dynamic()`, `omp_get_dynamic()`
    - Dynamic ~ whether # of threads in a block can be changed when the block is running



# Runtime Library

- Locks

Synchronization pragmas and atomic operations are better

- Regular and nested locks

`omp_lock_t, omp_nest_lock_t`

- Initialization and destruction

`omp_init_lock(), omp_init_nest_lock()`

`omp_destroy_lock(), omp_destroy_nest_lock()`

- Acquiring (blocking) and releasing

`omp_set_lock(), omp_unset_lock(), ...`

- Acquiring the lock without blocking

`omp_test_lock(), omp_test_nest_lock()`



# Environmental Variables

- Environmental Variables
  - May affect the application without recompilation
  - **OMP\_NUM\_THREADS**
    - Max number of threads during execution
  - **OMP\_SCHEDULE**
    - Scheduling strategy for for-loop construct
    - The loop must have the strategy set to **runtime**
  - **OMP\_DYNAMIC**
    - Enables dynamic adjustment of threads in a block



# Thread Team Configuration

- Actual number of threads in parallel region
  1. `if(condition)` clause is evaluated
  2. `num_threads(n)` clause is used if present
  3. `omp_set_num_threads()` value is used
  4. `OMP_NUM_THREADS` env. value is used
  5. System default (typically # of CPU cores)
- Affinity
  - Whether threads are bound to CPU cores
  - `proc_bind(policy)` clause
  - `OMP_PROC_BIND, OMP_PLACES`



# Nested Parallelism

- Nested Parallel Blocks

- Implementations may not support it
- Must be explicitly enabled

`omp_set_nested(), OMP_NESTED`

- Nesting depth may be limited

`omp_set_max_active_levels(),  
OMP_MAX_ACTIVE_LEVELS`

- More complex to get the right thread ID

`omp_get_level(), omp_get_active_level(),  
omp_get_ancestor_thread_num()`



# SIMD Support

- SIMD Instructions

- Generated by compiler when possible
- Hints may be provided by pragmas to loops  
`#pragma omp simd, #pragma omp for simd`

- Important clauses

- **safelen** – max. safe loop unroll width
- **simdlen** – recommended loop unroll width
- **aligned** – declaration about array(s) data alignment
- **linear** – declaration of variables with linear relation to iteration parameter

Use both SIMD  
and parallel for





# Accelerator Support

- Offload Support
  - Execute code on accelerator (GPU, FPGA, ...)  
`#pragma omp target ...`
  - Slightly more complicated
    - Memory has to be allocated on target device
      - And data transferred there and back
    - Or memory has to be mapped to target device
  - Target device may have more complex thread structure
    - OpenMP introduce thread **teams** directive