

Spark

Jakub Yaghob



Spark




- Apache Spark
 - Unified analytics engine for large-scale data processing
 - Initial release: 2014
 - Speed
 - Much faster than Hadoop
 - Easy of use
 - Write applications in Java, Python, R, and Scala
 - Interactive shell
 - Generality
 - Combine streaming, ML, SQL, and analytics
 - Runs everywhere
 - Spark runs on Hadoop, Mesos, Kubernetes, standalone, or in a cloud



Spark

- Contains
 - High-level API in Java, Python, R, and Scala
 - Optimized engine for general execution graph
 - DAG
 - MapReduce with only 2 levels
 - High-level tools
 - SparkSQL
 - MLlib
 - Spark Streaming
 - GraphX






Spark application

- Components
 - Application
 - Independent set of processes on a cluster (+isolation, -no shared data among applications without writing to an external storage)
 - Coordinated by **SparkContext** object in your main program (driver program)
 - Driver program must listen for incoming connections from its executors – must be network addressable from executors
 - Driver schedules tasks – it should be on the same local network
 - Driver program has a web UI – tasks, executors, storage
 - Cluster manager
 - Allocate resources across applications
 - Executor
 - On nodes of the cluster
 - Executes computations and stores data
 - Stays up for the duration of the application
 - Runs task in multiple threads



Spark application structure

- App structure
 - Create **SparkContext** object
 - Connect to the cluster manager(s)
 - Acquire executors from cluster manager(s)
 - Send your application code to the executors
 - Pass Java/Python code/files to the **SparkContext**
 - **SparkContext** sends tasks to the executors to run





Cluster managers

- Cluster managers
 - Spark is agnostic to the underlying cluster manager
 - Supported cluster managers
 - Standalone
 - Simple CM included with Spark, easy to setup a cluster
 - Apache Mesos
 - General CM able to run Hadoop MapReduce and service applications
 - Hadoop YARN
 - Resource manager in Hadoop 2
 - Kubernetes
 - Automated deployment, scaling, and management of containerized applications



Data holders

- Possible data holders
 - RDD
 - Resilient distributed dataset
 - Immutable collection of elements partitioned across the nodes that can be operated in parallel
 - Possibly in-memory
 - Automatically recover from node failures
 - Core API, from initial release, all languages
 - Dataset
 - Distributed collection of data
 - Since v1.6, available only for Java and Scala
 - Strong typing and lambdas from RDD
 - Using Spark SQL optimized execution engine
 - DataFrame
 - Like Dataset, organized into named columns
 - Since v1.3, all languages
 - Like a table in a relational DB



DataFrame

- Construction
 - Well known data file formats
 - CSV, JSON, ...
 - Other sources
 - External DB, existing RDD, tables from Hadoop, ...
- Transformations
 - Create a new DataFrame from the existing one
 - Lazily evaluated, triggered by an action
- Actions
 - Returns a result to the driver or writes to disk

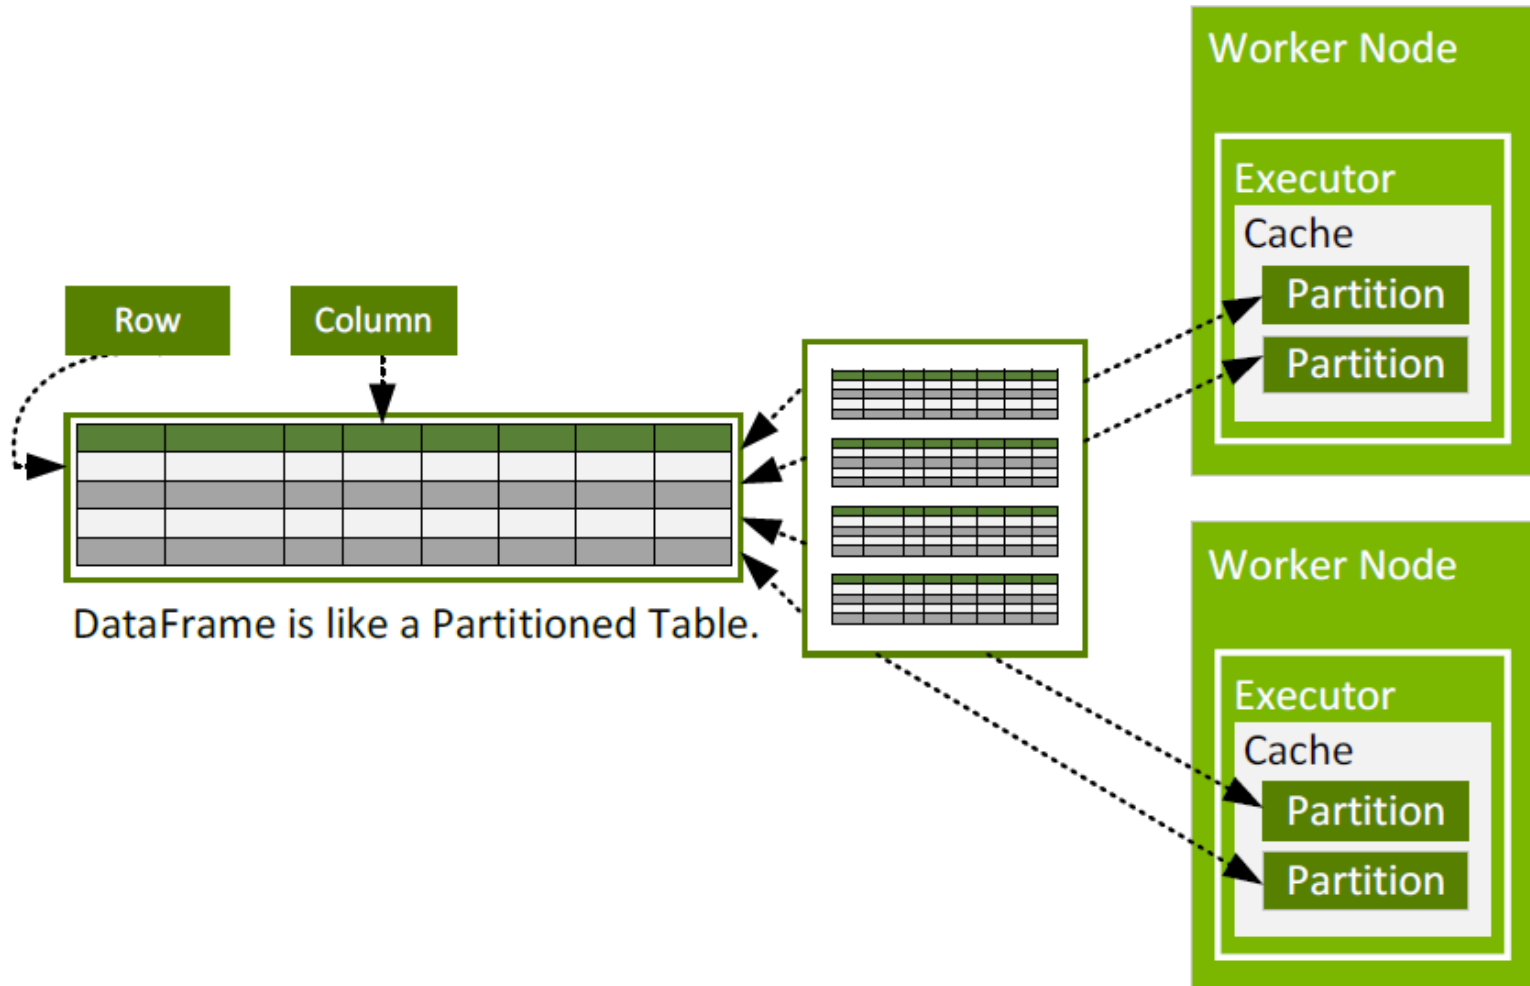
DataFrame transformations and actions



- Transformations (examples)
 - select – select a set of columns
 - join – join with another DataFrame
 - groupBy – groups using specified columns
 - filter – filter rows using a condition (bool or string)
- Actions (examples)
 - show(n) – display the first n rows
 - count – number of rows in DataFrame
 - collect – return data back to the driver
- Example
 - `df.groupBy("hour").count().show(4)`
 - `df.filter(df.age > 3).collect()`



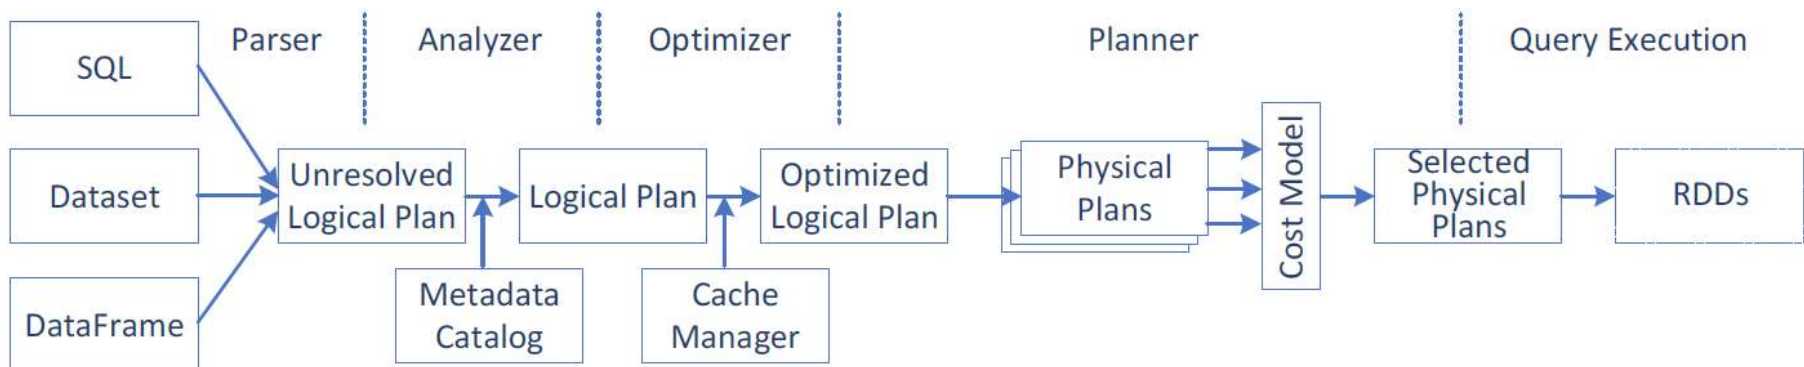
DataFrame partitioning





Spark query execution

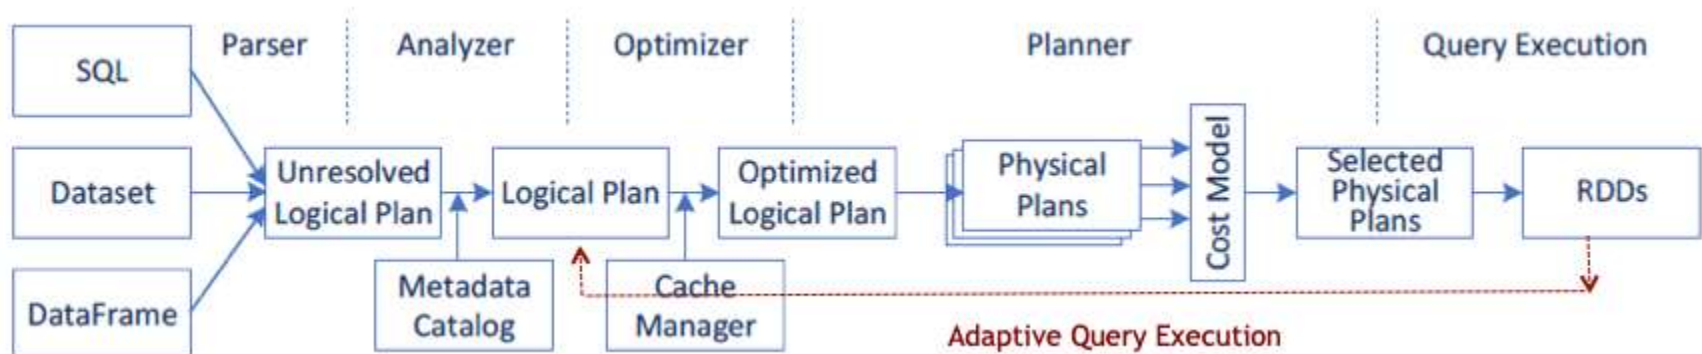
- Query execution steps
 - Create a logical plan
 - Transform the logical plan to a physical plan
 - Generate code
 - Execute the tasks on a cluster





Advanced Spark 3.x features

- Adaptive Query Execution (AQE)
 - Runtime statistics retrieved from completed stages of the query plan are used to re-optimize the execution plan
 - Dynamically coalescing shuffle partitions
 - Combine adjacent small partitions into bigger partitions, reducing the number of tasks
 - Dynamically switching join strategies
 - Optimize join strategy at runtime based on the join relation size
 - Dynamically optimizing skew joins
 - Detect data skew and split skew partitions into smaller sub-partitions





Advanced Spark 3.x features

- Dynamic partition pruning
 - Data warehouse queries
 - One or more fact tables referencing any number of dimensional tables
 - Pruning at runtime by reusing the dimension table broadcast results in hash joins
- Accelerator-aware scheduling
 - GPU/CUDA with RAPIDS



Parallel aspects

- **SparkContext** type
 - `parallelize(col, slices)`
 - Distribute a local collection to form an RDD
 - `accumulator(ival)`
 - Creates an Accumulator with initial value
 - `broadcast(ival)`
 - Broadcast read-only Broadcast variable to the cluster
- **RDD** type
 - `aggregate(zeroval, seqop, combop)`
 - Aggregate elements of each partition and then the results for all partitions
 - `barrier`
 - All tasks launched together
 - `cache`
 - Partitions cached in memory
 - `reduce(op)`



Shared variables

- Accumulator
 - Worker tasks can call `add(v)` operation
 - Only driver can read accumulator by calling `value()` function
 - No other operations are defined
- Broadcast
 - Cached read-only variable
 - Tasks can read it by calling `value()` function



Launching application

- Interactive shell
 - Only Scala and Python
 - `YOUR_SPARK_HOME/bin/pyspark`
- Standalone applications
 - Unified launcher
 - `YOUR_SPARK_HOME/bin/spark-submit`
 - Important parameters
 - `--master` – URL of the master node for the cluster
 - `--class` – entry point for the application
 - Master URL
 - `local` - locally with one worker thread (no para)
 - `local[K]` – locally with K worker threads
 - `local[*]` – locally with max number of worker threads (=cores)
 - `spark://HOST:PORT` – standalone Spark cluster master
 - `CM://HOST:PORT` – connect to cluster manager [mesos,yarn,k8s]

Launching application in SLURM environment



- Use prepared environment
 - Environment home
 - `/home/_teaching/para/04-spark`
 - Spark cluster startup script
 - `spark-slurm.sh`
 - Requires
 - Spark Charliecloud image directory – `spark`
 - Network interface with IP networking – `eno1` for w[201-208]
 - R/W directory – your home or project dir, mounted as `/mnt/1`
 - Application – path from the container (`/mnt/1/...`)
 - Launch the script using `sbatch` command