



GPGPU with SYCL

NPRG042: Programming in Parallel Environment

Martin Kruliš

GPU vs. CPU



- CPU

- Few cores per chip
- General purpose cores
- Processing different threads
- Huge caches to reduce memory latency
 - Locality of reference problem



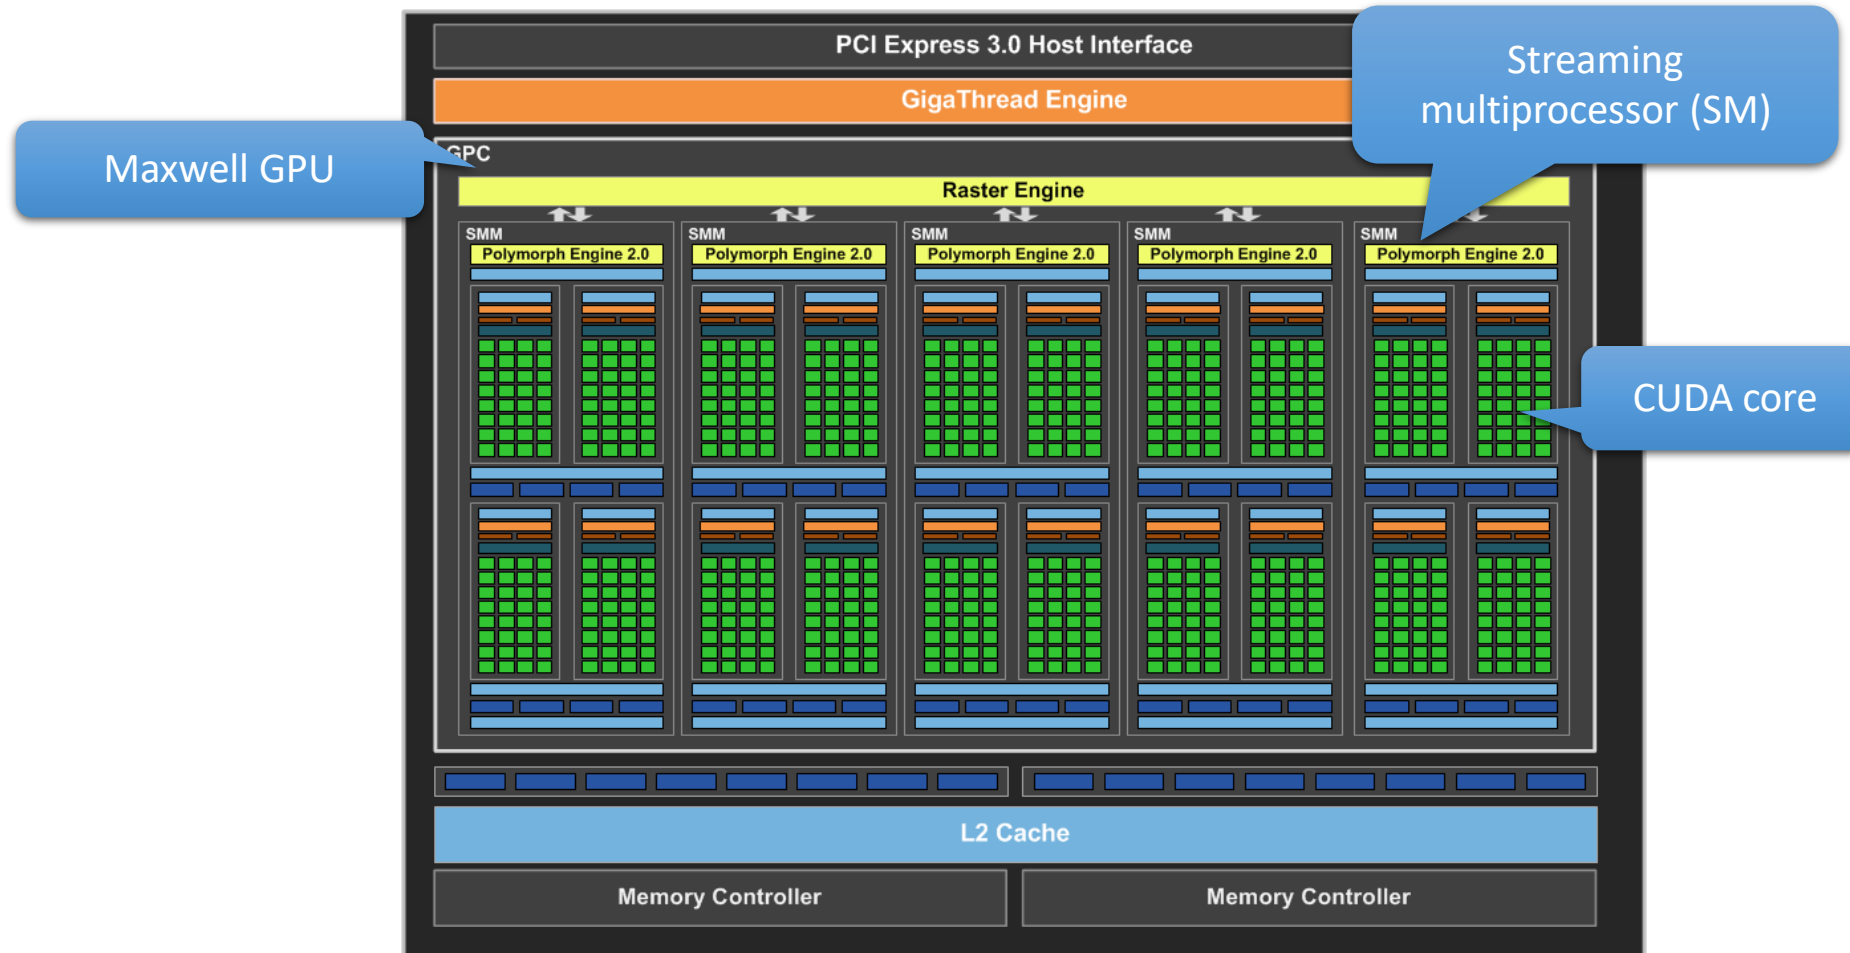
- GPU

- Many cores per chip
- Cores specialized for numeric computations
- SIMT thread processing
- Huge amount of threads and fast context switch
 - Results in more complex memory transfers



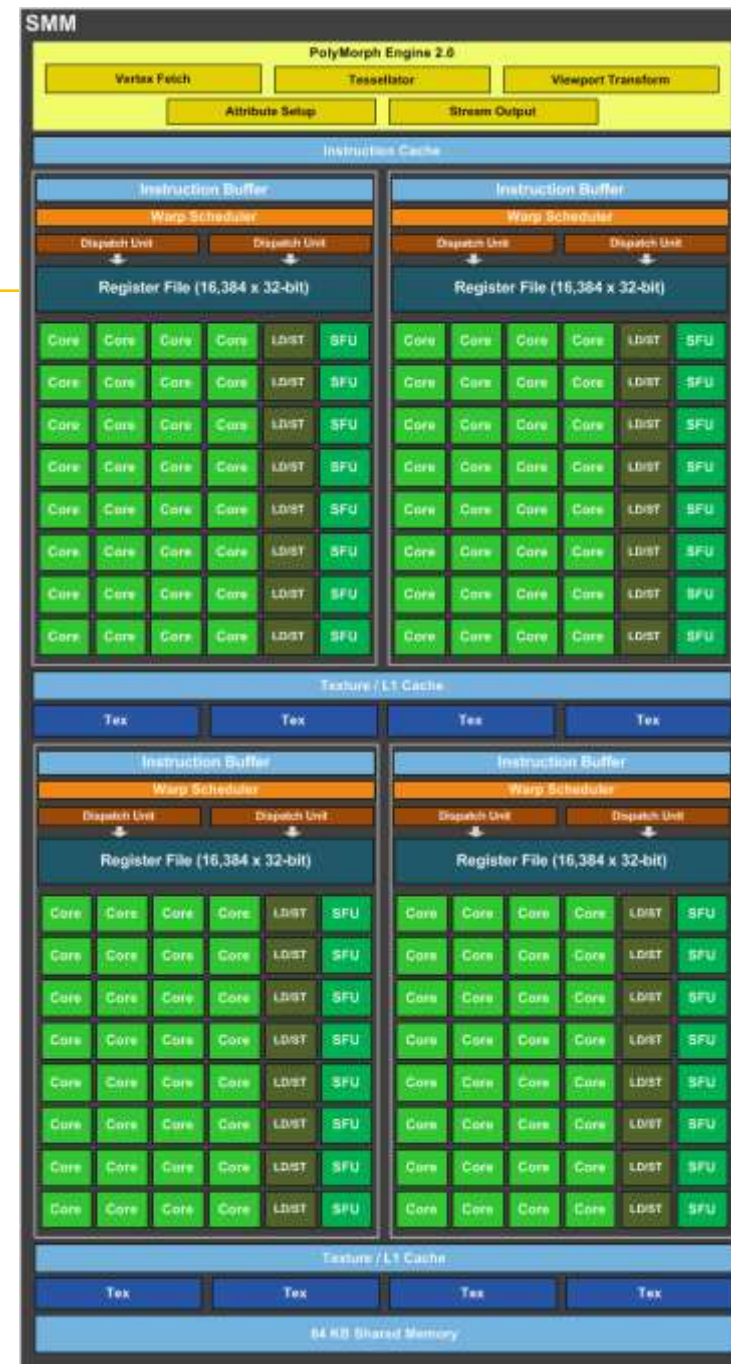
Architecture
Convergence

GPU Architecture



GPU Architecture

- Maxwell Architecture
 - 4 identical parts
 - 32 cores
 - 64 kB shared memory
 - 2 instruction schedulers
 - CC 5.0
 - SMM
(Streaming Multiprocessor - Maxwell)



GPU Architecture

- Volta SM
 - 7.x CC
 - 8x tensor core
 - (64 FMA/clock each)





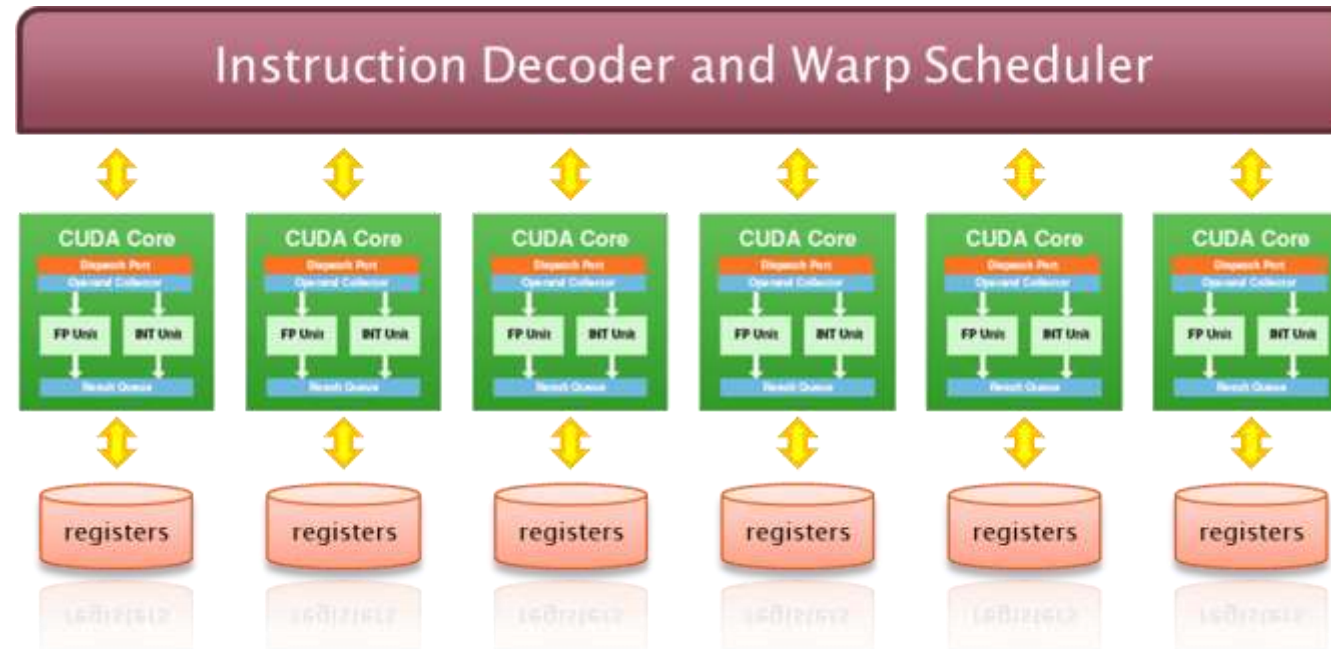
Execution Model

- Data Parallelism
 - Many data elements are processed concurrently by the same routine
 - GPUs are designed under this particular paradigm
 - Also have limited ways to express task parallelism
- Threading Execution Model
 - One function (kernel) is executed in many threads (work items)
 - Much more lightweight than the CPU threads
 - Threads are grouped into blocks (work groups) of the same size



Single Instruction Multiple Threads

- Single Instruction Multiple Threads
 - All cores are executing the same instruction
 - Each core has its own set of registers



SIMT vs. SIMD



- Single Instruction Multiple Threads

- Width-independent programming model
- Serial-like code
- Achieved by hardware with a little help from the compiler
- Allows code divergence

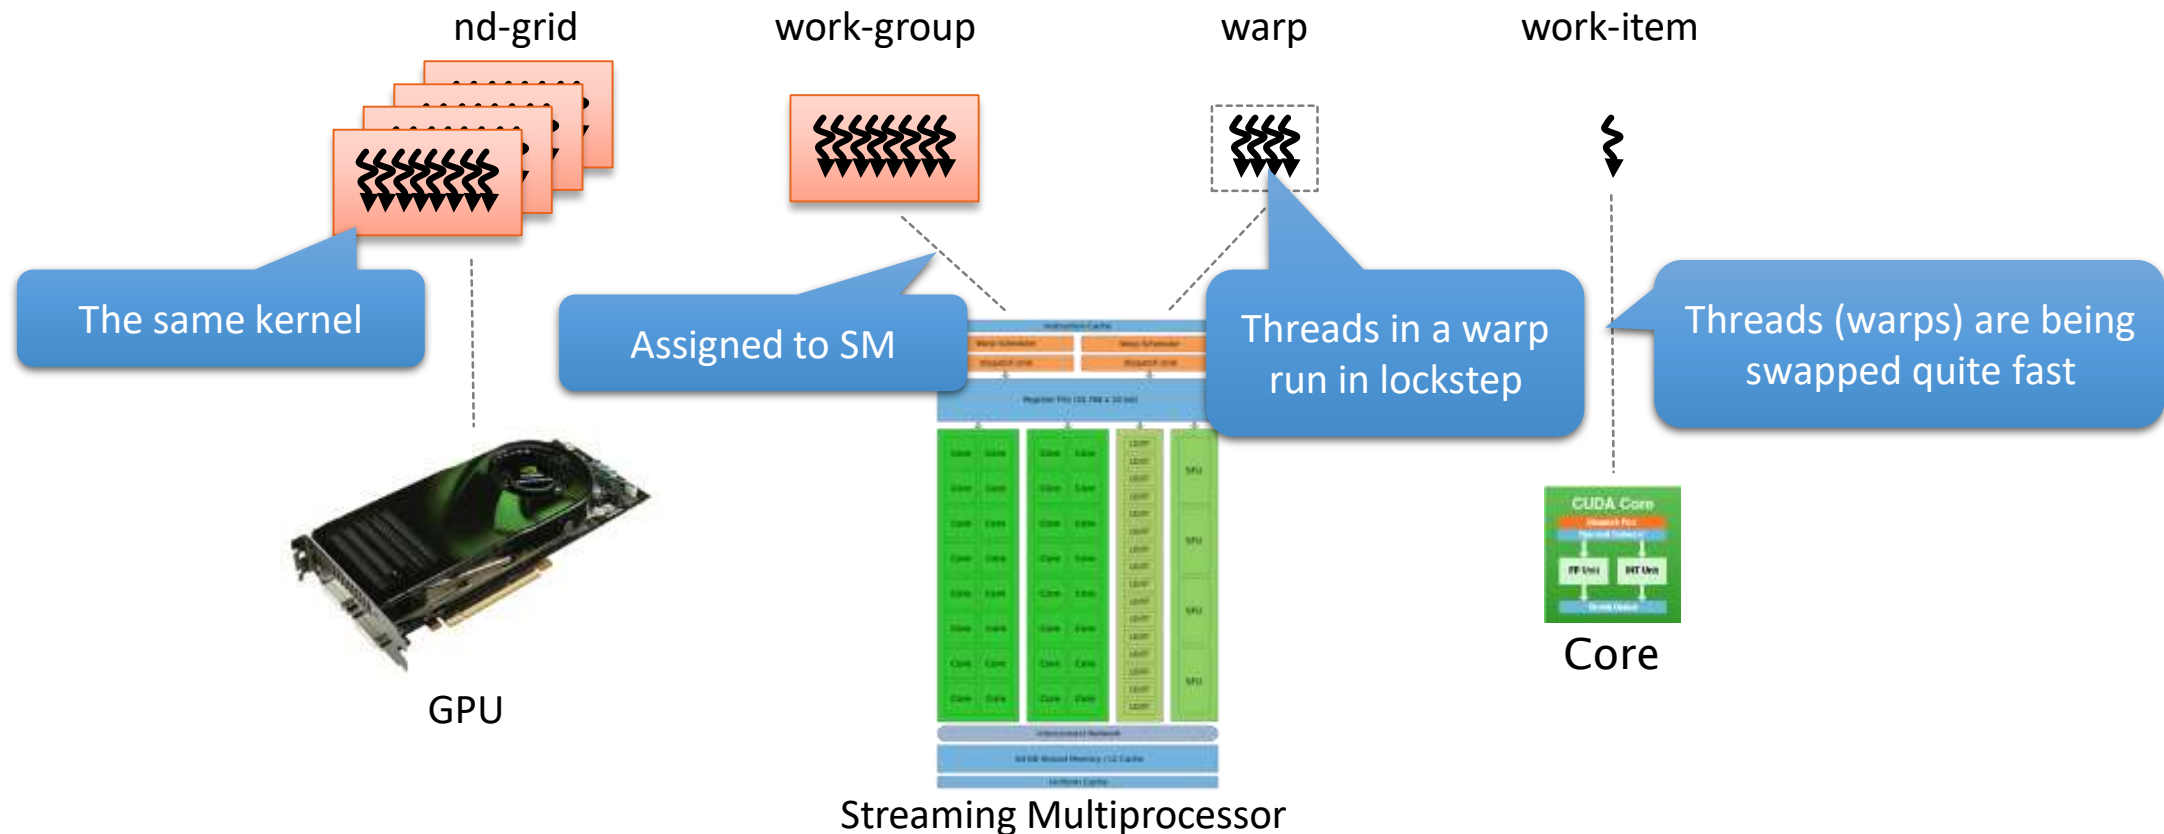
- Single Instruction Multiple Data

- Explicitly expose the width of the SIMD vector
- Special instructions
- Generated by the compiler or directly written by a programmer
- Code divergence is usually not supported or tedious



Workload Distribution

- Nd-grid and work elements assignment



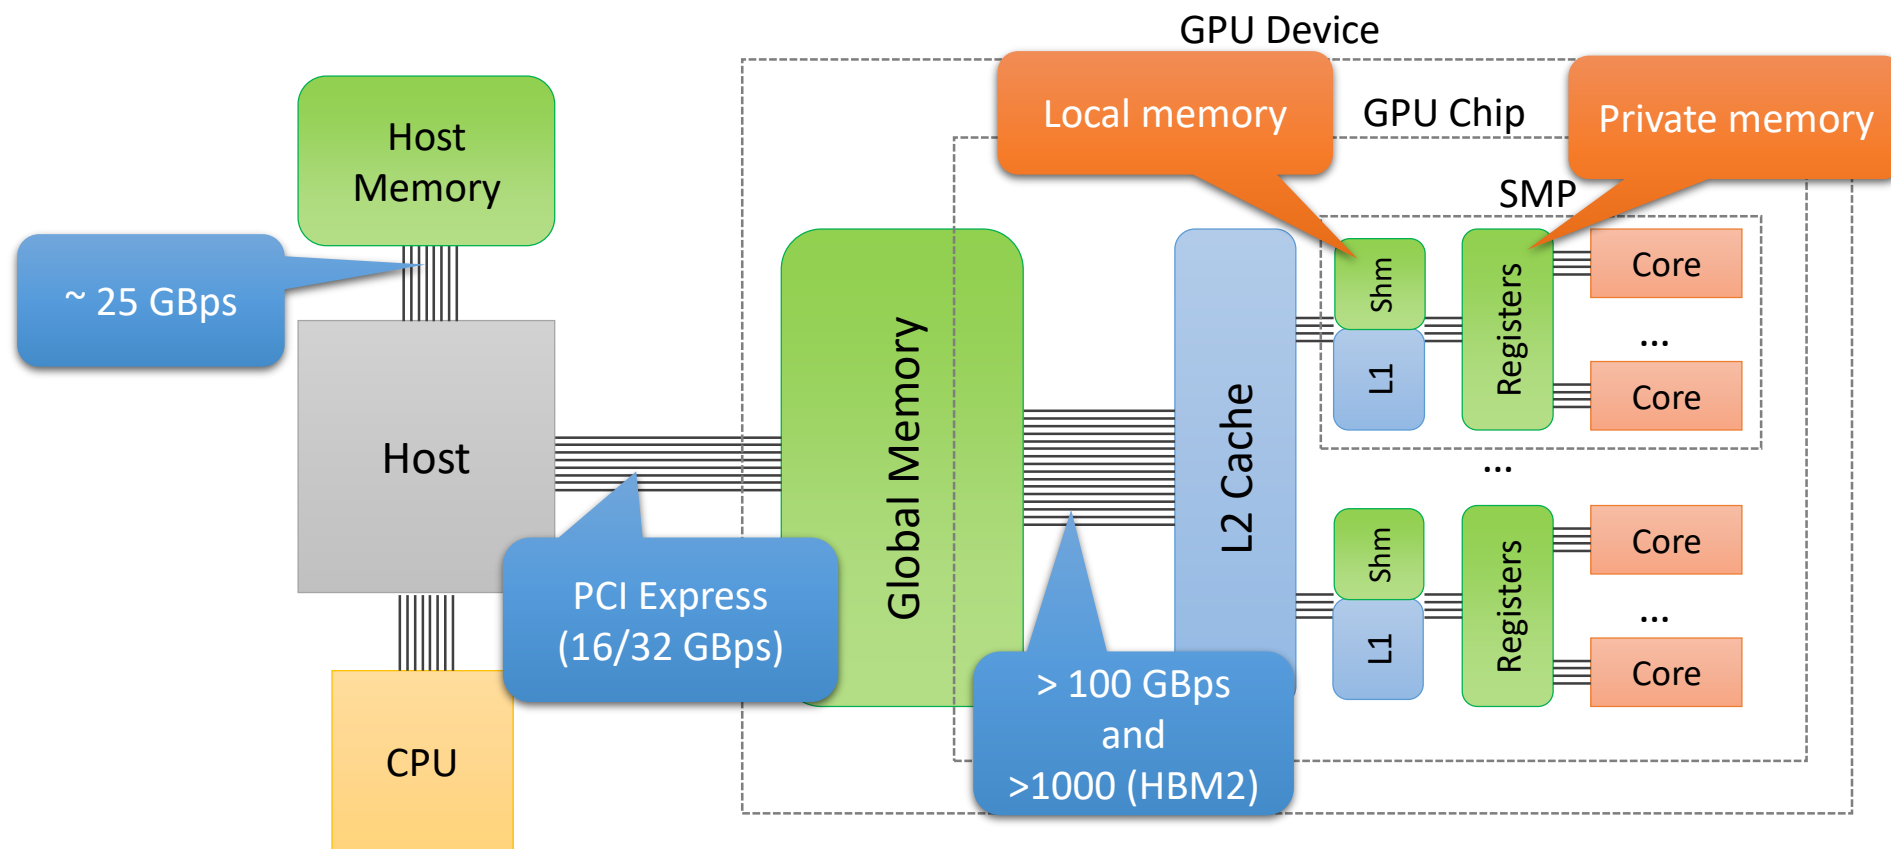
Lockstep Execution



- Lockstep
 - Multiple threads (warp) execute the same instruction simultaneously
 - Simplifies the design of the SMs
 - The programmer needs to know
 - When optimizing the code
 - Not expressed explicitly in SYCL
 - Cannot rely on that (e.g., for sync.)
 - Divergence is possible
 - But not very efficient
 - Volta significantly loosens lockstep
 - Threads may diverge in hardware



Memory



Memory



- Host-device transfers
 - Implicit (declared by accessors, handled by the scheduler) – no control
 - Explicit `handler.copy(src, dest);`
- Registers
 - All local variables of the kernel
 - Limited size, registry spilling (cached in L1)
- Local (shared) memory
 - Special memory shared among work-group
 - Very fast, used for (manual) caching or data exchange (items in group)
`auto shm = sycl::accessor<int, 1, sycl::access::target::local> (...);`

Code Example



```
static auto exh =
    [] (sycl::exception_list l)
{
    for (std::exception_ptr const &e : l) {
        try {
            std::rethrow_exception(e);
        }
        catch (std::exception const &e) {
            std::terminate();
        }
    }
};

std::vector<int> A, B, C;
auto selector = sycl::gpu_selector_v;

try {
    sycl::buffer a_buf(A);
    sycl::buffer b_buf(B);
```

```
sycl::range<1> rng{ A.size() };
sycl::buffer c_buf(C.data(), rng);

sycl::queue q(selector, exh);
q.submit([&] (sycl::handler &h) {
    sycl::accessor a(a_buf, h, sycl::read_only);
    sycl::accessor b(b_buf, h, sycl::read_only);
    sycl::accessor c(c_buf, h, sycl::write_only,
        sycl::no_init);

    h.parallel_for(rng, [=] (auto i) {
        c[i] = a[i] + b[i];
    });
});
q.wait_and_throw();

} catch (std::exception const &e) {
    // print e and terminate
}
```



Synchronization (Reminder)

- Host side

Queue can be configured as in-order

- `queue.wait()` ;
- Events
 - Returned by queue operations
`auto event = queue.submit(...)` ;
 - `event.wait()` ;
 - Events can be passed as dependencies to some operations
`queue.submit(..., event, ...)`

Kernel execution overhead is in the range of μ s

- In kernel

- Barriers
 - `nd_item::barrier()`
 - **Only within the work group!**
- Atomic operations
 - Similar to C++
`auto v =`
`sycl::ext::oneapi::atomic_ref`
`<int, order, scope, space>(a[0]) ;`
`v += something;`

Available both for local and global memory, quite efficient (if no collisions)


Lockstep Implementation



- Masking Instructions

- In case of data-driven branches
 - if-else conditions, while loops, ...
- All branches are traversed, threads mask their execution in invalid branches

```
if (threadIdx.x % 2 == 0) {  
    ... even threads code ...  
} else {  
    ... odd threads code ...  
}
```






Reducing Thread Divergence

- Work Reorganization

- In case the workload is imbalanced
- “Cheap” load balancing can lead to better occupancy

The overhead must not eliminate the gain



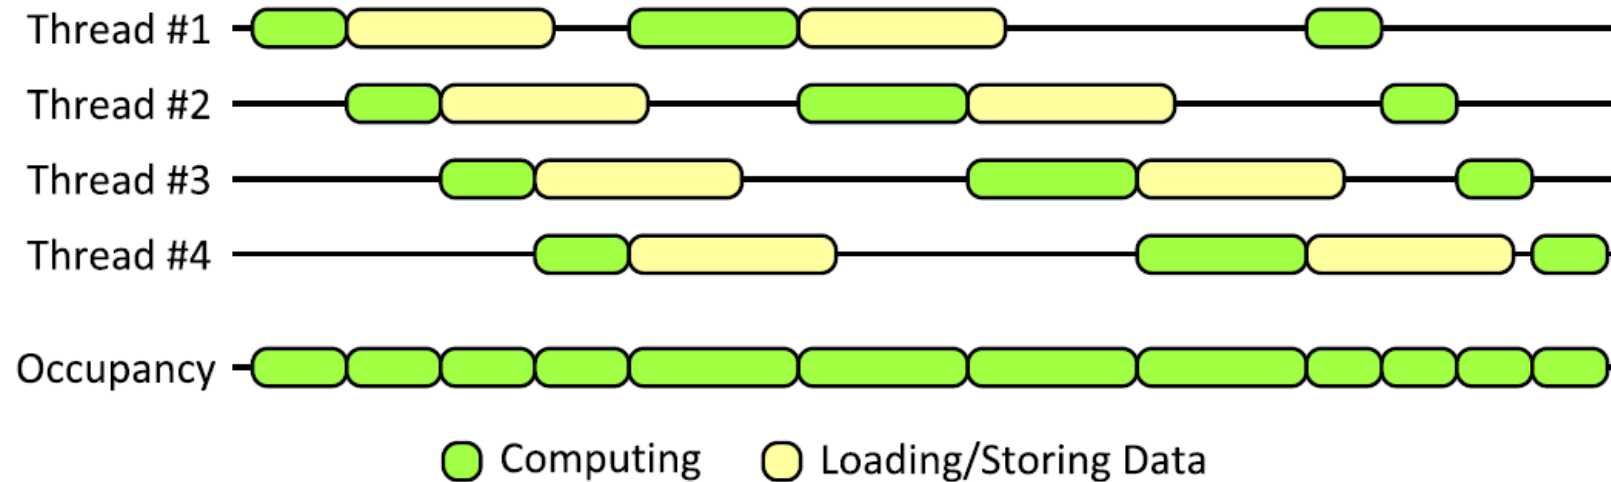
- Example

- Matrix with dimensions not divisible by warp size
- Item (i, j) has linear index $i * \text{width} + j$



Optimizations: Hiding Latency

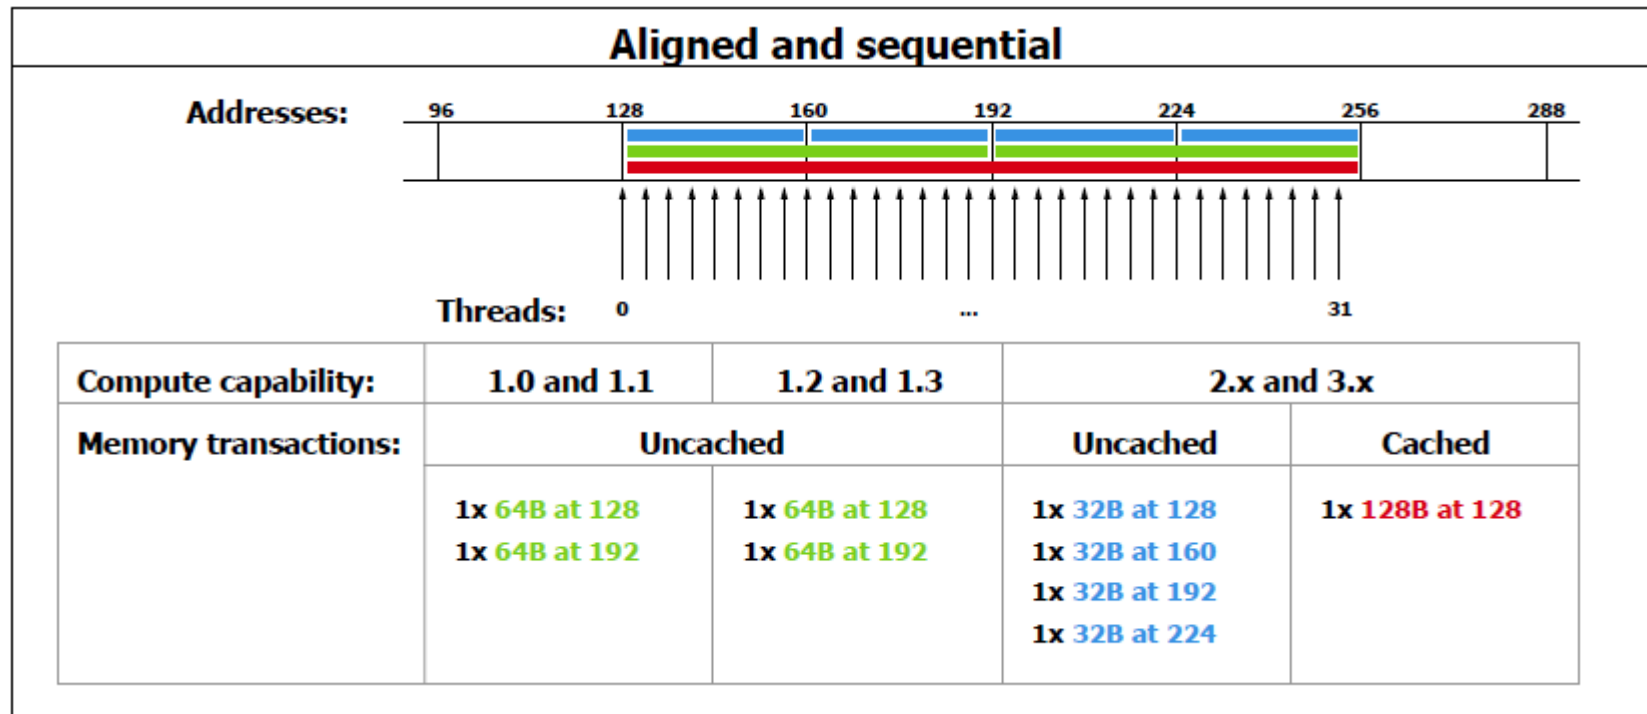
- Fast Context Switch
 - When a warp gets stalled
 - E.g., by data load/store
 - Scheduler switch to next active warp





Optimizations: Global Memory


- Access Patterns
 - Perfectly aligned sequential access





Optimizations: Global Memory

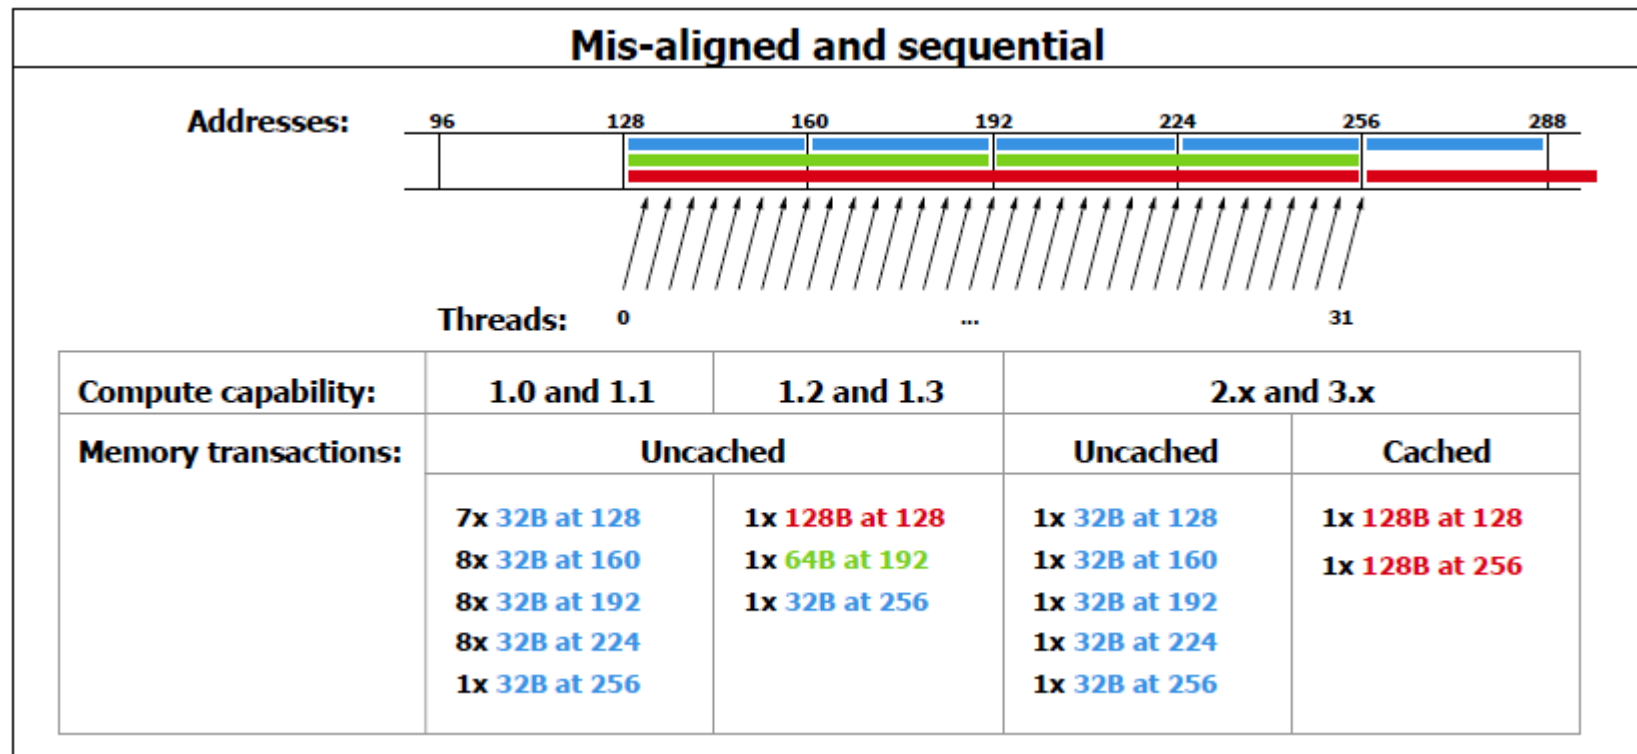
- Access Patterns
 - Perfectly aligned with permutation





Optimizations: Global Memory


- Access Patterns
 - Continuous sequential, but misaligned



Optimizations: Global Memory




- Coalesced Loads Impact



Optimizations: Local (Shared) Memory




- Linear Addressing
 - Each thread in warp access different memory bank
 - No collisions



Optimizations: Local (Shared) Memory




- Linear Addressing with Stride
 - Each thread access $2 \cdot i$ -th item
 - 2-way conflicts (2x slowdown) on CC < 3.0
 - No collisions on CC 3.x
 - Due to 64-bits per cycle throughput



Optimizations: Local (Shared) Memory




- Linear Addressing with Stride
 - Each thread access $3 \cdot i$ -th item
 - No collisions, since the number of banks is not divisible by the stride



Optimizations: Local (Shared) Memory




- Broadcast
 - One set of threads access value in bank #12 and the remaining threads access value in bank #20
 - Broadcasts simultaneously





Optimizations: Host-device Transfers


- Making a Good Use of Overlapping
 - Split the work into smaller fragments
 - Create a pipeline effect (load, process, store)



Heterogeneous Programming



- GPU
 - “Independent” device
 - Controlled by host
 - Used for “offloading”
- Host Code
 - Needs to be designed in a way that
 - Utilizes GPU(s) efficiently
 - Utilize CPU while GPU is working
 - CPU and GPU do not wait for each other



Discussion

